

a5

May 8, 2025

# 1 Dados

## 1.1 Caminhos

```
[1]: datasets_path      = '/home/jose-roberto/Documents/Disciplinas/INF492/assignment/  
    ↪a5/datasets'  
models_path           = '/home/jose-roberto/Documents/Disciplinas/INF492/assignment/  
    ↪a5/models/'  
tensorboard_path      = '/home/jose-roberto/Documents/Disciplinas/INF492/assignment/  
    ↪a5/train_runs/'
```

## 1.2 Dataloader

```
[2]: from torch.utils.data import DataLoader  
import torchvision  
  
import matplotlib.pyplot as plt  
import numpy as np  
  
def my_imshow(img, dataset, numImages=10):  
  
    if dataset == 'cifar10':  
        img = img / 2 + 0.5      # unnormalize  
  
    img = torchvision.utils.make_grid(img[:numImages],nrow=numImages//2)  
  
    npimg = img.numpy()  
    npimg = np.transpose(npimg, (1, 2, 0))  
  
    plt.axis('off')  
    plt.imshow(npimg)  
    plt.show()  
  
def show_images(train_loader, test_loader, dataset, numImages=10) :  
    print('Train samples')  
    # get some random training images  
    dataiter = iter(train_loader)  
    images = next(dataiter)[0]
```

```

my_imshow(images, dataset, numImages)

print('Test samples')
# get some random training images
dataiter = iter(test_loader)
images = next(dataiter)[0]
my_imshow(images, dataset, numImages)

def get_data_cifar10 ( batch_size , show_image=False, numImages=10 ) :

    my_transform = torchvision.transforms.Compose([
        torchvision.transforms.Resize(28),
        torchvision.transforms.ToTensor(),
        torchvision.transforms.Normalize(mean=[0.5],std=[0.
↪5])

    ])

    train_dataset = torchvision.datasets.CIFAR10(
        root=f'{datasets_path}/train/',
        train=True,
        transform=my_transform,
        download=False
    )
    test_dataset = torchvision.datasets.CIFAR10(
        root=f'{datasets_path}/test/',
        train=False,
        transform=my_transform,
        download=False
    )
    train_loader = DataLoader(train_dataset,
        batch_size=batch_size,
        shuffle=True
    )
    test_loader = DataLoader(test_dataset,
        batch_size=batch_size,
        shuffle=False
    )

    if show_image :
        show_images(train_loader, test_loader, 'cifar10', numImages)

    return train_loader, test_loader, len(train_dataset)

def get_data_mnist ( batch_size , show_image=False, numImages=10 ) :

    train_dataset = torchvision.datasets.mnist.MNIST(

```

```

        root=f'{datasets_path}/train/',
        train=True,
        transform=torchvision.transforms.ToTensor(),
        download=False
    )
    test_dataset = torchvision.datasets.mnist.MNIST(
        root=f'{datasets_path}/test/',
        train=False,
        transform=torchvision.transforms.ToTensor(),
        download=False
    )

    train_loader = DataLoader(train_dataset, batch_size=batch_size,
↪shuffle=True)
    test_loader = DataLoader(test_dataset, batch_size=batch_size, shuffle=False)


    if show_image :
        show_images(train_loader, test_loader, 'mnist', numImages)

    return train_loader, test_loader, len(train_dataset)


```

```
[3]: get_data_mnist(batch_size=256, show_image=True, numImages=16);
```

Train samples




Test samples




```
[4]: get_data_cifar10(batch_size=256, show_image=True, numImages=10);
```

Train samples



Test samples



## 2 Rede

### 2.1 Arquitetura

```
[24]: import torch
import torch.nn as nn
import torch.nn.functional as F

class LeNet(nn.Module) :

    def __init__(self, num_classes=10, n_channels=1, relu=False):
```

```

    super().__init__()

    self.conv1 = nn.Conv2d(in_channels=n_channels, out_channels=6,
↪kernel_size=5, padding=2)
    self.act1 = nn.Sigmoid() if not relu else nn.ReLU()
    self.pool1 = nn.AvgPool2d(kernel_size=2, stride=2)

    self.conv2 = nn.Conv2d(in_channels=6, out_channels=16, kernel_size=5)
    self.pool2 = nn.AvgPool2d(kernel_size=2, stride=2)

    out_channels_conv3 = 120 if n_channels == 1 else 480
    self.conv3 = nn.Conv2d(in_channels=16, out_channels=out_channels_conv3,
↪kernel_size=5)

    in_features_dense1 = out_channels_conv3
    self.dense1 = nn.Linear(in_features=in_features_dense1, out_features=84)
    self.act2 = nn.Tanh() if not relu else nn.ReLU()

    self.dense2 = nn.Linear(in_features=84, out_features=num_classes)

def forward(self, x, debug=False):
    if debug : print(f'Shape de entrada: {x.shape}')
    x = self.conv1(x)
    if debug : print(f'Shape de entrada: {x.shape}')
    x = self.act1(x)
    x = self.pool1(x)
    if debug : print(f'Shape de entrada: {x.shape}')
    x = self.conv2(x)
    if debug : print(f'Shape de entrada: {x.shape}')
    x = self.act1(x)
    x = self.pool2(x)
    if debug : print(f'Shape de entrada: {x.shape}')
    x = self.conv3(x)
    if debug : print(f'Shape de entrada: {x.shape}')
    x = self.act1(x)
    x = x.view(x.size(0), -1)
    if debug : print(f'Shape de entrada: {x.shape}')
    x = self.dense1(x)
    if debug : print(f'Shape de entrada: {x.shape}')
    x = self.act2(x)
    x = self.dense2(x)
    if debug : print(f'Shape de entrada: {x.shape}')
    x = self.act2(x)

    return x

```

## 2.2 Informações sobre a rede

```
[6]: import torch

if torch.cuda.is_available():
    my_device = torch.device("cuda:0")
else:
    my_device = torch.device("cpu")

print(f"Running on {my_device.type}.")

net = LeNet(num_classes=10, n_channels=1)
net = net.to(my_device)

a = torch.rand( (1, 1, 28, 28) )

b = net( a.to(my_device), debug=True)
```

Running on cpu.

```
Shape de entrada: torch.Size([1, 1, 28, 28])
Shape de entrada: torch.Size([1, 6, 28, 28])
Shape de entrada: torch.Size([1, 6, 14, 14])
Shape de entrada: torch.Size([1, 16, 10, 10])
Shape de entrada: torch.Size([1, 16, 5, 5])
Shape de entrada: torch.Size([1, 120, 1, 1])
Shape de entrada: torch.Size([1, 120])
Shape de entrada: torch.Size([1, 84])
Shape de entrada: torch.Size([1, 10])
```

```
[7]: from torchsummary import summary
summary(net, input_size=(1,28,28), batch_size=256)
del net, a, b
```

Layer (type)	Output Shape	Param #
Conv2d-1	[256, 6, 28, 28]	156
Sigmoid-2	[256, 6, 28, 28]	0
AvgPool2d-3	[256, 6, 14, 14]	0
Conv2d-4	[256, 16, 10, 10]	2,416
Sigmoid-5	[256, 16, 10, 10]	0
AvgPool2d-6	[256, 16, 5, 5]	0
Conv2d-7	[256, 120, 1, 1]	48,120
Sigmoid-8	[256, 120, 1, 1]	0
Linear-9	[256, 84]	10,164
Tanh-10	[256, 84]	0
Linear-11	[256, 10]	850
Tanh-12	[256, 10]	0

Total params: 61,706  
Trainable params: 61,706  
Non-trainable params: 0

-----  
Input size (MB): 0.77  
Forward/backward pass size (MB): 28.54  
Params size (MB): 0.24  
Estimated Total Size (MB): 29.54  
-----

```
[25]: net = LeNet(num_classes=10, n_channels=3)
net = net.to(my_device)

a = torch.rand( (1, 3, 28, 28) )

b = net( a.to(my_device), debug=True)

from torchsummary import summary
summary(net, input_size=(3,28,28), batch_size=256)
del net, a, b
```

Shape de entrada: torch.Size([1, 3, 28, 28])  
Shape de entrada: torch.Size([1, 6, 28, 28])  
Shape de entrada: torch.Size([1, 6, 14, 14])  
Shape de entrada: torch.Size([1, 16, 10, 10])  
Shape de entrada: torch.Size([1, 16, 5, 5])  
Shape de entrada: torch.Size([1, 480, 1, 1])  
Shape de entrada: torch.Size([1, 480])  
Shape de entrada: torch.Size([1, 84])  
Shape de entrada: torch.Size([1, 10])

-----

Layer (type)	Output Shape	Param #
Conv2d-1	[256, 6, 28, 28]	456
Sigmoid-2	[256, 6, 28, 28]	0
AvgPool2d-3	[256, 6, 14, 14]	0
Conv2d-4	[256, 16, 10, 10]	2,416
Sigmoid-5	[256, 16, 10, 10]	0
AvgPool2d-6	[256, 16, 5, 5]	0
Conv2d-7	[256, 480, 1, 1]	192,480
Sigmoid-8	[256, 480, 1, 1]	0
Linear-9	[256, 84]	40,404
Tanh-10	[256, 84]	0
Linear-11	[256, 10]	850
Tanh-12	[256, 10]	0

=====

Total params: 236,606  
Trainable params: 236,606

Non-trainable params: 0

-----  
Input size (MB): 2.30  
Forward/backward pass size (MB): 29.95  
Params size (MB): 0.90  
Estimated Total Size (MB): 33.14  
-----

## 2.3 Treinamento

```
[9]: from torch.utils.tensorboard import SummaryWriter

import torch.optim

import matplotlib.pyplot as plt
import numpy as np
from tqdm import tqdm

import copy

from datetime import datetime

def train ( dataset, epochs=100, lr=1e-1, prefix='', upper_bound=99.0,
    ↪device='cpu',
        relu=False, save=False, debug=False, plot_histograms=False,
    ↪lambda_reg=0) :

    if dataset == 'mnist' :
        batch_size = 256
        train_loader, test_loader, dataset_size = get_data_mnist(batch_size,
            ↪show_image=True
            )

        num_classes = 10
        n_channels = 1
    elif dataset == 'cifar10' :
        batch_size = 256
        train_loader, test_loader, dataset_size = get_data_cifar10(batch_size,
            ↪show_image=True
            )

        num_classes = 10
        n_channels = 3
    else :
        print('Dataset loader not implemented.')
        return None
```



```

net = LeNet( num_classes, n_channels, relu=relu )
net.to(device)

optimizer = torch.optim.SGD(net.parameters(), lr, weight_decay=lambda_reg)
loss = nn.CrossEntropyLoss()

now = datetime.now()
suffix = now.strftime("%Y%m%d_%H%M%S")
prefix = prefix + '-' + suffix if prefix != '' else suffix

writer = SummaryWriter( log_dir=tensorboard_path+prefix )

writer.add_graph(net, next(iter(train_loader))[0].to(my_device))

accuracies = []
max_accuracy = -1.0

for epoch in tqdm( range(epochs) , desc='Training epochs...' ) :
    net.train()
    for idx, (train_x, train_label) in enumerate(train_loader):
        train_x = train_x.to(device)
        train_label = train_label.to(device)

        optimizer.zero_grad()

        predict_y = net( train_x )

        # Loss:
        error = loss( predict_y , train_label.long() )

        writer.add_scalar( 'Loss/train', error,
                           idx+( epoch*(dataset_size//batch_size) ) )

        error.backward()
        optimizer.step()

        # Accuracy:
        predict_ys = torch.max( predict_y, axis=1 )[1]
        correct     = torch.sum(predict_ys == train_label)

        writer.add_scalar( 'Accuracy/train', correct/train_x.size(0),
                           idx+( epoch*(dataset_size//batch_size) ) )

        if debug and idx % 10 == 0 :
            print(f'idx: {idx}, _error: {error}')

    if plot_histograms :

```

```

        plot_histograms_tensorboard(writer, net, epoch)

    accuracy = validate(net, test_loader, device=device)
    accuracies.append(accuracy.cpu())
    writer.add_scalar( 'Accuracy/test', accuracy, epoch )

    if accuracy > max_accuracy:
        best_model = copy.deepcopy(net)
        max_accuracy = accuracy
        print(f'Saving Best Model with Accuracy: {max_accuracy:3.4f}')

    print( f'Epoch: {epoch+1:3d} | Accuracy : {accuracy:3.4f}%' )

    if accuracy > upper_bound :
        break

    if save :
        path = f'{models_path}-{ 'relu-' if relu else
↪ '' }LeNet5-{dataset}-{max_accuracy:.2f}.pkl'
        torch.save(best_model.state_dict(), path)
        print('Model saved in:',path)

plt.plot(accuracies)

writer.flush()
writer.close()

return best_model

```

```

[10]: def plot_histograms_tensorboard ( writer, net, epoch ) :
    writer.add_histogram('Bias/conv1',    net.conv1.bias,        epoch)
    writer.add_histogram('Weight/conv1',  net.conv1.weight,     epoch)
    writer.add_histogram('Grad/conv1',    net.conv1.weight.grad, epoch)

    writer.add_histogram('Bias/conv2',    net.conv2.bias,        epoch)
    writer.add_histogram('Weight/conv2',  net.conv2.weight,     epoch)
    writer.add_histogram('Grad/conv2',    net.conv2.weight.grad, epoch)

    writer.add_histogram('Bias/conv3',    net.conv3.bias,        epoch)
    writer.add_histogram('Weight/conv3',  net.conv3.weight,     epoch)
    writer.add_histogram('Grad/conv3',    net.conv3.weight.grad, epoch)

```

## 2.4 Validação

```
[11]: def validate ( model , data , device='cpu' ) :  
    model.eval()  
    correct = 0  
    sum = 0  
  
    for idx, (test_x, test_label) in enumerate(data) :  
        test_x = test_x.to(device)  
        test_label = test_label.to(device)  
        predict_y = model( test_x ).detach()  
        predict_ys = torch.max( predict_y, axis=1 )[1]  
        sum = sum + test_x.size(0)  
        correct = correct + torch.sum(predict_ys == test_label)  
  
    return correct*100./sum
```

## 3 Execução

### 3.1 Treina

```
[12]: if torch.cuda.is_available():  
    my_device = torch.device("cuda:0")  
else:  
    my_device = torch.device("cpu")  
print(f"Running on {my_device.type}.")
```

Running on cpu.

#### 3.1.1 MNIST - Sigmoid e Tahn

```
[38]: epochs = 30  
dataset = 'mnist' # 'cifar10'  
lr = 1.3e0  
lambda_reg = 0  
  
prefix = 'LeNet-{}-e-{}-lr-{}'.format(dataset, epochs, lr)  
  
net = train( dataset=dataset, epochs=epochs, lr=lr, prefix=prefix ,  
    ↪upper_bound=100, device=my_device,  
        relu=False, save=True, debug=False, plot_histograms=True,  
    ↪lambda_reg=lambda_reg )
```

Train samples



Test samples



Training epochs...: 3%| | 1/30 [00:15<07:29, 15.51s/it]

Saving Best Model with Accuracy: 11.3500

Epoch: 1 | Accuracy : 11.3500%

Training epochs...: 7%| | 2/30 [00:31<07:29, 16.07s/it]

Epoch: 2 | Accuracy : 11.3500%

Training epochs...: 10%| | 3/30 [00:48<07:17, 16.20s/it]

Epoch: 3 | Accuracy : 11.3500%

Training epochs...: 13%| | 4/30 [01:04<06:56, 16.03s/it]


Epoch: 4 | Accuracy : 11.3500%

Training epochs...: 17%| | 5/30 [01:20<06:43, 16.13s/it]

Epoch: 5 | Accuracy : 11.3500%

Training epochs...: 20%| | 6/30 [01:36<06:26, 16.09s/it]  
 Epoch: 6 | Accuracy : 11.3500%  
 Training epochs...: 23%| | 7/30 [01:52<06:07, 15.99s/it]  
 Epoch: 7 | Accuracy : 11.3500%  
 Training epochs...: 27%| | 8/30 [02:08<05:54, 16.11s/it]  
 Epoch: 8 | Accuracy : 11.3500%  
 Training epochs...: 30%| | 9/30 [02:25<05:40, 16.22s/it]  
 Epoch: 9 | Accuracy : 11.3500%  
 Training epochs...: 33%| | 10/30 [02:40<05:22, 16.11s/it]  
 Saving Best Model with Accuracy: 50.9600  
 Epoch: 10 | Accuracy : 50.9600%  
 Training epochs...: 37%| | 11/30 [02:56<05:04, 16.03s/it]  
 Saving Best Model with Accuracy: 86.6700  
 Epoch: 11 | Accuracy : 86.6700%  
 Training epochs...: 40%| | 12/30 [03:12<04:47, 15.95s/it]  
 Saving Best Model with Accuracy: 89.7200  
 Epoch: 12 | Accuracy : 89.7200%  
 Training epochs...: 43%| | 13/30 [03:29<04:34, 16.12s/it]  
 Saving Best Model with Accuracy: 92.8500  
 Epoch: 13 | Accuracy : 92.8500%  
 Training epochs...: 47%| | 14/30 [03:45<04:20, 16.26s/it]  
 Saving Best Model with Accuracy: 93.6000  
 Epoch: 14 | Accuracy : 93.6000%  
 Training epochs...: 50%| | 15/30 [04:02<04:07, 16.48s/it]  
 Saving Best Model with Accuracy: 94.9500  
 Epoch: 15 | Accuracy : 94.9500%  
 Training epochs...: 53%| | 16/30 [04:19<03:50, 16.50s/it]  
 Epoch: 16 | Accuracy : 93.9800%  
 Training epochs...: 57%| | 17/30 [04:36<03:37, 16.71s/it]  
 Epoch: 17 | Accuracy : 94.7500%  
 Training epochs...: 60%| | 18/30 [04:52<03:18, 16.53s/it]  
 Saving Best Model with Accuracy: 96.6400  
 Epoch: 18 | Accuracy : 96.6400%  
 Training epochs...: 63%| | 19/30 [05:08<02:59, 16.32s/it]

Saving Best Model with Accuracy: 97.0100  
 Epoch: 19 | Accuracy : 97.0100%  
 Training epochs...: 67%| | 20/30 [05:24<02:44, 16.42s/it]  
 Epoch: 20 | Accuracy : 96.6000%  
 Training epochs...: 70%| | 21/30 [05:40<02:26, 16.31s/it]  
 Saving Best Model with Accuracy: 97.4400  
 Epoch: 21 | Accuracy : 97.4400%  
 Training epochs...: 73%| | 22/30 [05:57<02:10, 16.36s/it]  
 Epoch: 22 | Accuracy : 96.2700%  
 Training epochs...: 77%| | 23/30 [06:14<01:55, 16.49s/it]  
 Saving Best Model with Accuracy: 97.6300  
 Epoch: 23 | Accuracy : 97.6300%  
 Training epochs...: 80%| | 24/30 [06:30<01:38, 16.46s/it]  
 Saving Best Model with Accuracy: 97.8100  
 Epoch: 24 | Accuracy : 97.8100%  
 Training epochs...: 83%| | 25/30 [06:47<01:22, 16.56s/it]  
 Epoch: 25 | Accuracy : 97.7800%  
 Training epochs...: 87%| | 26/30 [07:03<01:05, 16.41s/it]  
 Saving Best Model with Accuracy: 97.9100  
 Epoch: 26 | Accuracy : 97.9100%  
 Training epochs...: 90%| | 27/30 [07:19<00:49, 16.40s/it]  
 Epoch: 27 | Accuracy : 97.6200%  
 Training epochs...: 93%| | 28/30 [07:35<00:32, 16.26s/it]  
 Saving Best Model with Accuracy: 98.1900  
 Epoch: 28 | Accuracy : 98.1900%  
 Training epochs...: 97%| | 29/30 [07:51<00:16, 16.12s/it]  
 Epoch: 29 | Accuracy : 97.0900%  
 Training epochs...: 100%| | 30/30 [08:09<00:00, 16.30s/it]  
 Epoch: 30 | Accuracy : 98.1800%  
 Model saved in: /home/jose-  
 roberto/Documents/Disciplinas/INF492/assignment/a5/models/LeNet5-mnist-98.19.pkl




### 3.1.2 MNIST - ReLu

```
[40]: epochs = 30
dataset = 'mnist' # 'cifar10'
lr = 5e-2
lambda_reg = 0


prefix = 'relu-LeNet-{}-e-{}-lr-{}'.format(dataset, epochs, lr)

net_relu = train( dataset=dataset, epochs=epochs, lr=lr, prefix=prefix ,
    ↳upper_bound=100, device=my_device,
        relu=True, save=True, debug=False, plot_histograms=True,
    ↳lambda_reg=lambda_reg )
```

Train samples



Test samples



Training epochs...: 3%| | 1/30 [00:15<07:18, 15.12s/it]

Saving Best Model with Accuracy: 12.7600

Epoch: 1 | Accuracy : 12.7600%

Training epochs...: 7%| | 2/30 [00:29<06:56, 14.86s/it]

Saving Best Model with Accuracy: 49.2600

Epoch: 2 | Accuracy : 49.2600%

Training epochs...: 10%| | 3/30 [00:44<06:42, 14.91s/it]

Saving Best Model with Accuracy: 77.9200

Epoch: 3 | Accuracy : 77.9200%

Training epochs...: 13%| | 4/30 [01:00<06:36, 15.27s/it]


Saving Best Model with Accuracy: 83.6800

Epoch: 4 | Accuracy : 83.6800%



Training epochs...: 17%| | 5/30 [01:16<06:29, 15.56s/it]  
 Saving Best Model with Accuracy: 90.5100  
 Epoch: 5 | Accuracy : 90.5100%  
 Training epochs...: 20%| | 6/30 [01:33<06:21, 15.89s/it]  
 Saving Best Model with Accuracy: 93.0100  
 Epoch: 6 | Accuracy : 93.0100%  
 Training epochs...: 23%| | 7/30 [01:50<06:13, 16.26s/it]  
 Saving Best Model with Accuracy: 93.9400  
 Epoch: 7 | Accuracy : 93.9400%  
 Training epochs...: 27%| | 8/30 [02:05<05:54, 16.11s/it]  
 Saving Best Model with Accuracy: 94.4200  
 Epoch: 8 | Accuracy : 94.4200%  
 Training epochs...: 30%| | 9/30 [02:23<05:44, 16.42s/it]  
 Saving Best Model with Accuracy: 95.9400  
 Epoch: 9 | Accuracy : 95.9400%  
 Training epochs...: 33%| | 10/30 [02:39<05:28, 16.43s/it]  
 Saving Best Model with Accuracy: 96.3800  
 Epoch: 10 | Accuracy : 96.3800%  
 Training epochs...: 37%| | 11/30 [02:55<05:08, 16.24s/it]  
 Saving Best Model with Accuracy: 96.7100  
 Epoch: 11 | Accuracy : 96.7100%  
 Training epochs...: 40%| | 12/30 [03:11<04:53, 16.30s/it]  
 Saving Best Model with Accuracy: 97.0300  
 Epoch: 12 | Accuracy : 97.0300%  
 Training epochs...: 43%| | 13/30 [03:27<04:33, 16.07s/it]  
 Epoch: 13 | Accuracy : 96.8600%  
 Training epochs...: 47%| | 14/30 [03:43<04:16, 16.02s/it]  
 Epoch: 14 | Accuracy : 96.7400%  
 Training epochs...: 50%| | 15/30 [04:00<04:04, 16.33s/it]  
 Saving Best Model with Accuracy: 97.2800  
 Epoch: 15 | Accuracy : 97.2800%  
 Training epochs...: 53%| | 16/30 [04:15<03:45, 16.13s/it]  
 Epoch: 16 | Accuracy : 97.1100%  
 Training epochs...: 57%| | 17/30 [04:31<03:27, 15.99s/it]  
 Epoch: 17 | Accuracy : 97.0700%

Training epochs...: 60%| | 18/30 [04:47<03:11, 15.98s/it]  
 Epoch: 18 | Accuracy : 96.9800%  
 Training epochs...: 63%| | 19/30 [05:03<02:54, 15.89s/it]  
 Saving Best Model with Accuracy: 97.7200  
 Epoch: 19 | Accuracy : 97.7200%  
 Training epochs...: 67%| | 20/30 [05:18<02:38, 15.83s/it]  
 Epoch: 20 | Accuracy : 94.9600%  
 Training epochs...: 70%| | 21/30 [05:34<02:22, 15.85s/it]  
 Epoch: 21 | Accuracy : 97.1800%  
 Training epochs...: 73%| | 22/30 [05:51<02:08, 16.08s/it]  
 Saving Best Model with Accuracy: 97.8100  
 Epoch: 22 | Accuracy : 97.8100%  
 Training epochs...: 77%| | 23/30 [06:08<01:54, 16.31s/it]  
 Epoch: 23 | Accuracy : 96.6900%  
 Training epochs...: 80%| | 24/30 [06:24<01:38, 16.38s/it]  
 Epoch: 24 | Accuracy : 97.6100%  
 Training epochs...: 83%| | 25/30 [06:41<01:21, 16.34s/it]  
 Saving Best Model with Accuracy: 98.0800  
 Epoch: 25 | Accuracy : 98.0800%  
 Training epochs...: 87%| | 26/30 [06:57<01:05, 16.25s/it]  
 Epoch: 26 | Accuracy : 97.8600%  
 Training epochs...: 90%| | 27/30 [07:13<00:48, 16.17s/it]  
 Epoch: 27 | Accuracy : 97.7700%  
 Training epochs...: 93%| | 28/30 [07:29<00:32, 16.15s/it]  
 Epoch: 28 | Accuracy : 97.8000%  
 Training epochs...: 97%| | 29/30 [07:45<00:16, 16.17s/it]  
 Saving Best Model with Accuracy: 98.2200  
 Epoch: 29 | Accuracy : 98.2200%  
 Training epochs...: 100%| | 30/30 [08:01<00:00, 16.04s/it]  
 Saving Best Model with Accuracy: 98.4000  
 Epoch: 30 | Accuracy : 98.4000%  
 Model saved in: /home/jose-  
 roberto/Documents/Disiplinas/INF492/assignment/a5/models/relu-  
 LeNet5-mnist-98.40.pkl




### 3.1.3 CIFAR10 - ReLu

```
[41]: epochs = 30
dataset = 'cifar10' # 'mnist'
lr = 5e-2
lambda_reg = 0


prefix = 'relu-LeNet-{}-e-{}-lr-{}'.format(dataset, epochs, lr)

net_relu = train( dataset=dataset, epochs=epochs, lr=lr, prefix=prefix ,
    ↳upper_bound=100, device=my_device,
        relu=True, save=True, debug=False, plot_histograms=True,
    ↳lambda_reg=lambda_reg )
```

Train samples



Test samples



Training epochs...: 3%| | 1/30 [00:22<10:43, 22.19s/it]

Saving Best Model with Accuracy: 16.6100

Epoch: 1 | Accuracy : 16.6100%

Training epochs...: 7%| | 2/30 [00:45<10:31, 22.56s/it]

Saving Best Model with Accuracy: 18.1400

Epoch: 2 | Accuracy : 18.1400%

Training epochs...: 10%| | 3/30 [01:09<10:34, 23.50s/it]

Saving Best Model with Accuracy: 24.7400

Epoch: 3 | Accuracy : 24.7400%


Training epochs...: 13%| | 4/30 [01:35<10:35, 24.46s/it]

Saving Best Model with Accuracy: 32.8100

Epoch: 4 | Accuracy : 32.8100%

Training epochs...: 17%| | 5/30 [02:00<10:17, 24.68s/it]  
 Saving Best Model with Accuracy: 38.0400  
 Epoch: 5 | Accuracy : 38.0400%  
 Training epochs...: 20%| | 6/30 [02:26<10:02, 25.12s/it]  
 Saving Best Model with Accuracy: 38.6900  
 Epoch: 6 | Accuracy : 38.6900%  
 Training epochs...: 23%| | 7/30 [02:52<09:46, 25.49s/it]  
 Saving Best Model with Accuracy: 41.9100  
 Epoch: 7 | Accuracy : 41.9100%  
 Training epochs...: 27%| | 8/30 [03:17<09:17, 25.34s/it]  
 Saving Best Model with Accuracy: 46.7100  
 Epoch: 8 | Accuracy : 46.7100%  
 Training epochs...: 30%| | 9/30 [03:43<08:51, 25.32s/it]  
 Saving Best Model with Accuracy: 47.0900  
 Epoch: 9 | Accuracy : 47.0900%  
 Training epochs...: 33%| | 10/30 [04:08<08:25, 25.27s/it]  
 Saving Best Model with Accuracy: 47.1300  
 Epoch: 10 | Accuracy : 47.1300%  
 Training epochs...: 37%| | 11/30 [04:34<08:03, 25.45s/it]  
 Epoch: 11 | Accuracy : 41.9400%  
 Training epochs...: 40%| | 12/30 [04:59<07:37, 25.43s/it]  
 Saving Best Model with Accuracy: 49.9500  
 Epoch: 12 | Accuracy : 49.9500%  
 Training epochs...: 43%| | 13/30 [05:25<07:12, 25.46s/it]  
 Saving Best Model with Accuracy: 50.0100  
 Epoch: 13 | Accuracy : 50.0100%  
 Training epochs...: 47%| | 14/30 [05:51<06:49, 25.61s/it]  
 Saving Best Model with Accuracy: 50.2100  
 Epoch: 14 | Accuracy : 50.2100%  
 Training epochs...: 50%| | 15/30 [06:16<06:24, 25.61s/it]  
 Epoch: 15 | Accuracy : 49.6300%  
 Training epochs...: 53%| | 16/30 [06:43<06:01, 25.84s/it]  
 Saving Best Model with Accuracy: 51.6200  
 Epoch: 16 | Accuracy : 51.6200%  
 Training epochs...: 57%| | 17/30 [07:08<05:34, 25.70s/it]

Epoch: 17 | Accuracy : 43.8900%  
 Training epochs...: 60%| | 18/30 [07:33<05:07, 25.65s/it]  
 Epoch: 18 | Accuracy : 46.8400%  
 Training epochs...: 63%| | 19/30 [07:59<04:42, 25.64s/it]  
 Epoch: 19 | Accuracy : 49.9800%  
 Training epochs...: 67%| | 20/30 [08:26<04:19, 25.94s/it]  
 Saving Best Model with Accuracy: 52.4700  
 Epoch: 20 | Accuracy : 52.4700%  
 Training epochs...: 70%| | 21/30 [08:52<03:55, 26.19s/it]  
 Epoch: 21 | Accuracy : 47.4000%  
 Training epochs...: 73%| | 22/30 [09:19<03:30, 26.33s/it]  
 Saving Best Model with Accuracy: 53.7300  
 Epoch: 22 | Accuracy : 53.7300%  
 Training epochs...: 77%| | 23/30 [09:45<03:02, 26.12s/it]  
 Epoch: 23 | Accuracy : 53.1400%  
 Training epochs...: 80%| | 24/30 [10:12<02:38, 26.36s/it]  
 Saving Best Model with Accuracy: 55.2500  
 Epoch: 24 | Accuracy : 55.2500%  
 Training epochs...: 83%| | 25/30 [10:37<02:10, 26.15s/it]  
 Saving Best Model with Accuracy: 55.4700  
 Epoch: 25 | Accuracy : 55.4700%  
 Training epochs...: 87%| | 26/30 [11:04<01:45, 26.40s/it]  
 Epoch: 26 | Accuracy : 53.9100%  
 Training epochs...: 90%| | 27/30 [11:31<01:19, 26.34s/it]  
 Epoch: 27 | Accuracy : 53.8900%  
 Training epochs...: 93%| | 28/30 [11:57<00:52, 26.37s/it]  
 Epoch: 28 | Accuracy : 53.9700%  
 Training epochs...: 97%| | 29/30 [12:24<00:26, 26.51s/it]  
 Epoch: 29 | Accuracy : 54.5500%  
 Training epochs...: 100%| | 30/30 [12:50<00:00, 25.68s/it]  
 Saving Best Model with Accuracy: 56.4200  
 Epoch: 30 | Accuracy : 56.4200%  
 Model saved in: /home/jose-  
 roberto/Documents/Disciplinas/INF492/assignment/a5/models/relu-  
 LeNet5-cifar10-56.42.pkl



## 4 Carregar Rede de arquivo

```
[43]: path = '/home/jose-roberto/Documents/Disciplinas/INF492/assignment/a5/models/
↳ LeNet5-mnist-98.19.pkl'
path_relu = '/home/jose-roberto/Documents/Disciplinas/INF492/assignment/a5/
↳ models/relu-LeNet5-mnist-98.40.pkl'
path_cifar10 = '/home/jose-roberto/Documents/Disciplinas/INF492/assignment/a5/
↳ models/relu-LeNet5-cifar10-56.42.pkl'

n_channels = 1

def load_LeNet ( device , path, n_channels=1, relu=False ) :
    net = LeNet(num_classes=10, n_channels=n_channels, relu=relu)
    net = net.to(device)
    net.load_state_dict(torch.load(path))
    net.eval()
    return net

net_mnist = load_LeNet(my_device, path)
```

```
net_relu = load_LeNet(my_device, path_relu, relu=True)
net_cifar10 = load_LeNet(my_device, path_cifar10, n_channels=3, relu=True)
```

## 5 Carregar dado do MNIST e inferir

```
[44]: import PIL
import torchvision
import numpy as np

def sample_and_predict_mnist ( net, seed=None, relu=False ) :

    if seed is not None :
        np.random.seed(seed)

    dataset = torchvision.datasets.MNIST(
                                                root=f'{datasets_path}/test/',
                                                train=False
                                            )

    i = np.random.randint(len(dataset))
    sample = dataset[i][0]

    x = torchvision.transforms.ToTensor()(sample).float()

    x = x.unsqueeze_(0)

    x = x.to(my_device)

    y = net ( x )

    if relu:
        y = torch.softmax(y, dim=1)

    confidence = torch.max(y, 1)[0]
    prediction = torch.max(y, 1)[1]

    print( 'Sample: {}'.format(i) )
    plt.axis('off')
    plt.imshow( sample , cmap='gray')

    confidence = confidence.data.cpu().numpy()[0]
    prediction = prediction.data.cpu().numpy()[0]

    return prediction, confidence, dataset[i][1]
```




```
[45]: prediction, confidence, label = sample_and_predict_mnist(net_mnist)
print( f'\nPredicted clas: {prediction} \nClassifier confidence:␣
↪{confidence*100:4.2f}% \nTrue label: {label}' )
```

Sample: 8333

Predicted clas: 6

Classifier confidence: 99.99%

True label: 6



```
[46]: prediction, confidence, label = sample_and_predict_mnist(net_relu, relu=True)
print( f'\nPredicted clas: {prediction} \nClassifier confidence:␣
↪{confidence*100:4.2f}% \nTrue label: {label}' )
```

Sample: 9209

Predicted clas: 2

Classifier confidence: 99.94%

True label: 2



## 6 Carregar dado do CIFAR10 e inferir

```
[47]: import PIL
import torchvision
import numpy as np

def sample_and_predict_cifar10 ( net, seed=None, relu=False ) :

    if seed is not None :
        np.random.seed(seed)

    my_transform = torchvision.transforms.Compose([
        torchvision.transforms.Resize(28),
        torchvision.transforms.ToTensor(),
        torchvision.transforms.Normalize(mean=[0.5],std=[0.
↪5])

    ])

    dataset = torchvision.datasets.CIFAR10(
        root=f'{datasets_path}/test/',
        transform=my_transform,
```

```

                                train=False
                                )

    i = np.random.randint(len(dataset))
    x = dataset[i][0]

    sample = np.array(x)
    sample = np.transpose(sample, (1, 2, 0))

    x = x.unsqueeze_(0)

    x = x.to(my_device)

    y = net ( x )

    if relu:
        y = torch.softmax(y, dim=1)

    confidence = torch.max(y, 1)[0]
    prediction = torch.max(y, 1)[1]

    print( 'Sample: {}'.format(i) )
    plt.axis('off')
    plt.imshow( sample )

    confidence = confidence.data.cpu().numpy()[0]
    prediction = prediction.data.cpu().numpy()[0]

    return prediction, confidence, dataset[i][1]

```

```

[48]: prediction, confidence, label = sample_and_predict_cifar10(net_cifar10,
    ↪relu=True)
print( f'\nPredicted clas: {prediction} \nClassifier confidence:
    ↪{confidence*100:4.2f}% \nTrue label: {label}' )

```

/tmp/ipykernel\_12544/2333663424.py:26: DeprecationWarning: \_\_array\_\_ implementation doesn't accept a copy keyword, so passing copy=False failed. \_\_array\_\_ must implement 'dtype' and 'copy' keyword arguments. To learn more, see the migration guide  
[https://numpy.org/devdocs/numpy\\_2\\_0\\_migration\\_guide.html#adapting-to-changes-in-the-copy-keyword](https://numpy.org/devdocs/numpy_2_0_migration_guide.html#adapting-to-changes-in-the-copy-keyword)

```


    sample = np.array(x)
Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers). Got range [-0.99215686..0.9843137].

```

Sample: 9810

Predicted clas: 5

Classifier confidence: 41.94%  
True label: 8



## 7 Conclusões

Analisando os gráficos plotados no TensorBoard, temos que a Loss para Sigmoid/Tahn estabiliza-se mais rapidamente, enquanto para ReLu há um decaimento maior. Para a acurácia, vemos que com ReLu há uma convergência muito mais rápida, enquanto para Sigmoid/Tahn os valores ficam estagnados e aumentam lentamente até se estabilizarem em uma faixa superior. Acompanhando o treino, vemos que para Sigmoid/Tahn foram necessárias 28 épocas para atingir 98% de acurácia, enquanto para ReLu foram necessárias 25.

E sobre a aplicação da mesma arquitetura no dataset CIFAR10, nota-se que o resultado máximo não foi satisfatório (aprox. 56% de acurácia). Isto era esperado, pois o dado é bem mais complexo comparado com o MNIST, contendo objetos complexos do mundo real, representados em três canais (RGB) e com bastante variações entre si.

## TensorBoard

TIME SERIES

SCALARS

INACTIVE

☐ Show data download links☐ Ignore outliers in chart scalingTooltip sorting  
method:

default

Smoothing



0.6



Horizontal Axis

STEP

RELATIVE

WALL

Runs

Write a regex to filter runs

LeNet-mnist-e-30-lr-1.3-2025  
0508\_202812relu-LeNet-mnist-e-30-lr-0.0  
5-20250508\_204535relu-LeNet-cifar10-e-30-lr-0.0  
5-20250508\_205420

TOGGLE ALL RUNS

/home/jose-roberto/Documents/  
Disciplinas/INF492/assignment/a5/  
train\_runs


Filter tags (regular expressions supported)

Accuracy

2 ^

test


tag: Accuracy/test



run to download

train

tag: Accuracy/train



run to download

Loss





## TensorBoard

TIME SERIES

SCALARS

INACTIVE

☐ Show data download links☐ Ignore outliers in chart scalingTooltip sorting  
method: default ▾

Smoothing

Horizontal Axis

STEP

RELATIVE

WALL

Runs

Write a regex to filter runs

☐ ☐ LeNet-mnist-e-30-lr-1.3-2025  
0508\_202812☐ ☐ relu-LeNet-mnist-e-30-lr-0.0  
5-20250508\_204535☐ ☐ relu-LeNet-cifar10-e-30-lr-0.0  
5-20250508\_205420

TOGGLE ALL RUNS

/home/jose-roberto/Documents/  
Disciplinas/INF492/assignment/a5/  
train\_runs

Filter tags (regular expressions supported)

Accuracy


2 ▾

Loss

^

train

tag: Loss/train



run to download ▾

